



Expectify

Rich Polymorphic Error Handling with `llvm::Expected<T>`

Stefan Gränitz

Freelance Dev C++ / Compilers / Audio

C++ User Group Berlin

19. September 2017

No answer for that..

YES! Exceptions

No Exceptions



Error Handling in the exception-free codebase

Whats the matter?

- ad-hoc approaches to indicate errors
- return `bool`, `int`, `nullptr` or `std::error_code`
- no concept for context information
- made for enumerable errors
- suffer from lack of enforcement

C++ has an answer for this

- ask the user to indicate error
- return `bool`, `int`, `nullptr` or `std::error_code`
- no concept for context information
- made multiple errors
- suffer from error enforcement

Exceptions

type-safe
handlers

user-defined
error types

total
enforcement

How to get these benefits without using exceptions?

```
Error foo(...);  
Expected<T> bar(...);
```

Polymorphic Error as a Return Value scheme

Idiomatic usage

```
Error foo(...);
```

```
// conversion to bool "checks" error
```

```
if (auto err = foo(...))  
    return err; // error case
```

```
// success case
```

Idiomatic usage

```
Error foo(...);  
Expected<T> bar(...);
```

```
foo(...); // unchecked Error triggers abort  
bar(...); // so does unchecked Expected
```



strong
enforcement

Idiomatic usage

```
Error foo(...);  
Expected<T> bar(...);
```



Don't silently
disappear or duplicate
(like Exceptions)

```
// errors can only be moved, not copied  
Error err1 = foo(...);  
Error err2 = std::move(err1);  
  
// ... same for Expected ...
```


Interface

```
class ErrorInfoBase {
public:
    virtual ~ErrorInfoBase() = default;

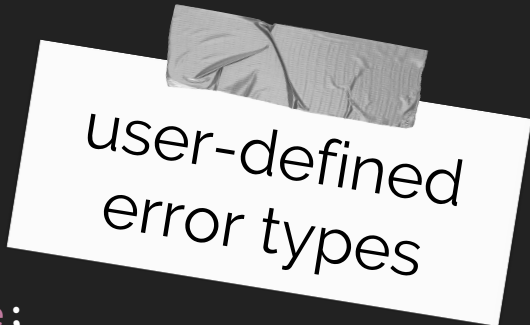
    /// Print an error message to an output stream.
    virtual void log(std::ostream &OS) const = 0;

    /// Return the error message as a string.
    virtual std::string message() const;

    /// Convert this error to a std::error_code.
    virtual std::error_code convertToErrorCode() const = 0;
};
```

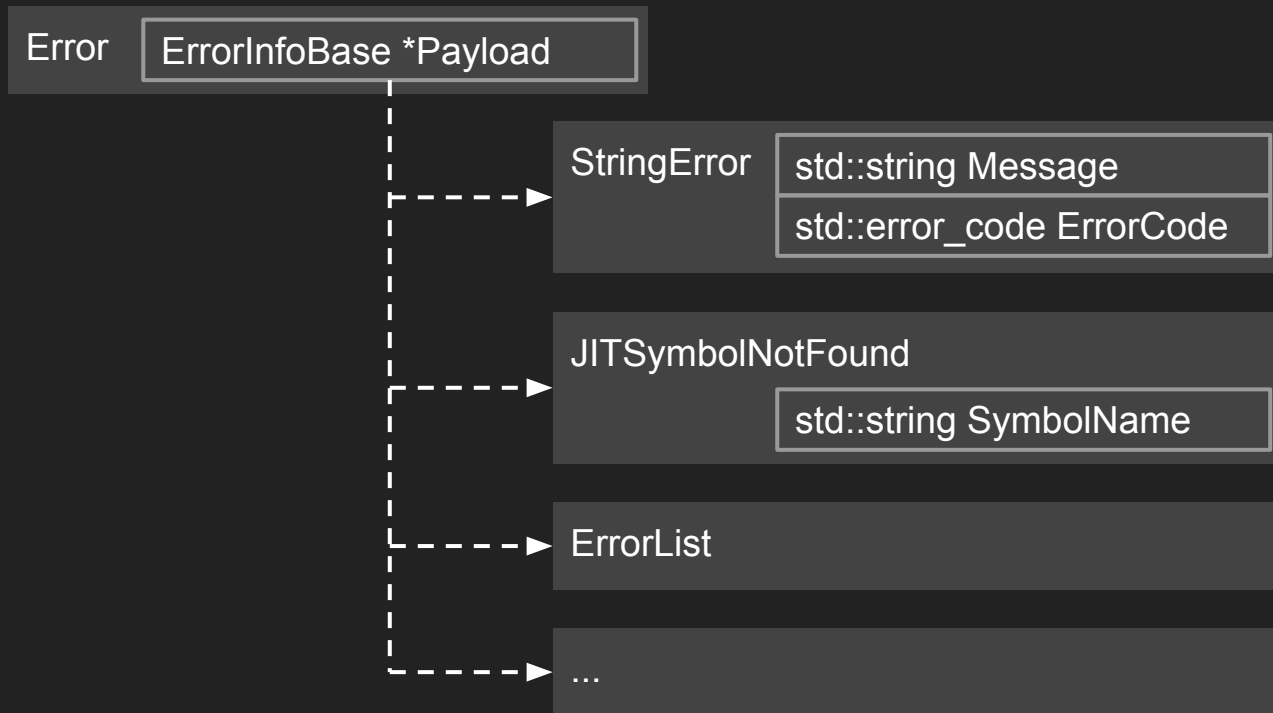
Implementation

```
class StringError : public ErrorInfo<StringError> {  
public:  
    static char ID;  
  
    StringError(std::string Msg, std::error_code EC);  
  
    void log(std::ostream &OS) const override;  
    std::error_code convertToErrorCode() const override;  
  
    const std::string &getMessage() const { return Msg; }  
  
private:  
    std::string Msg;  
    std::error_code EC;  
};
```

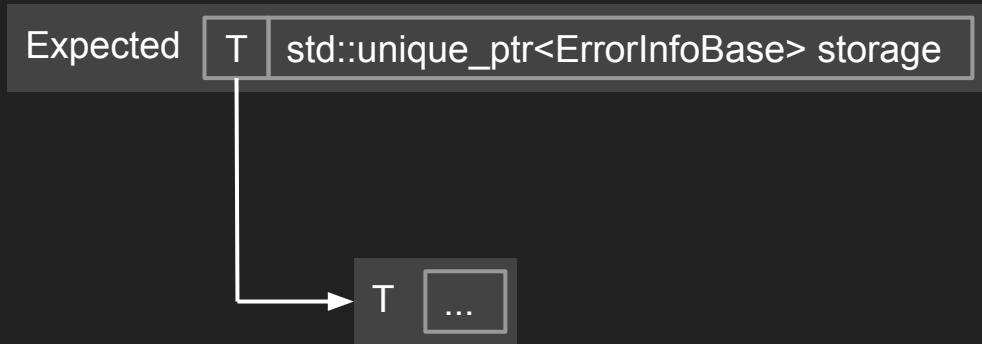


user-defined
error types

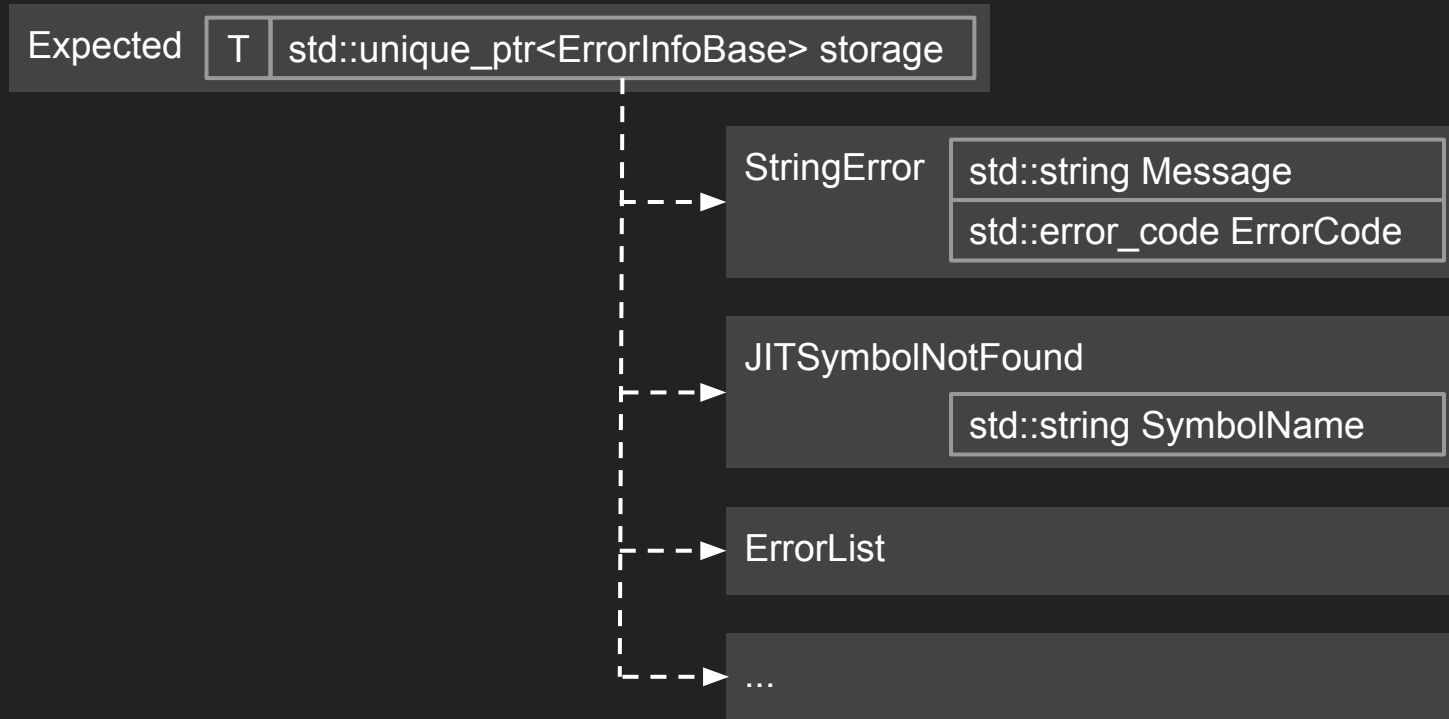
Composition



Composition



Composition



Utilities: make_error<T>

```
make_error<StringError>(
    "Bad executable",
    std::make_error_code(
        std::errc::executable_format_error));
```

Utilities: type-safe handlers

```
Error foo(...);
```

```
handleErrors(  
    foo(...),  
    [](const MyError &err){ ... },  
    [](SomeOtherError &err){ ... });
```



type-safe
handlers

Interop with `std::error_code`

```
std::error_code  errorToErrorCode(Error err);  
Error  errorCodeToError(std::error_code ec);
```

- useful when for porting a codebase
- similar to *Exploding Return Codes*

https://groups.google.com/forum/#!msg/comp.lang.c++.moderated/BkZqPfoq3ys/H_PMR8Sat4oJ

Example

```
bool simpleExample();  
  
int main() {  
    if (simpleExample())  
        // ... more code ...  
  
    return 0;  
}
```

Example . expected

```
bool simpleExample() {
    std::string fileName = "[a*.txt";
    Expected<GlobPattern> pattern = GlobPattern::create(std::move(fileName));

    if (auto err = pattern.takeError()) {
        logAllUnhandledErrors(std::move(err), std::cerr, "[Glob Error] ");
        return false;
    }

    return pattern->match("...");
}
```

Output: [Glob Error] invalid glob pattern: [a*.txt

Example . error_code

```
bool simpleExample() {
    std::string fileName = "[a*.txt";
    GlobPattern pattern;

    if (std::error_code ec = GlobPattern::create(fileName, pattern)) {
        std::cerr << "[Glob Error] " << getErrorDescription(ec) << ": ";
        std::cerr << fileName << "\n";
        return false;
    }

    return pattern.match("...");
}
```

Output: [Glob Error] invalid_argument: [a*.txt

Example . modified

```
std::error_code simpleExample(bool &result, std::string &errorFileName) {  
    GlobPattern pattern;  
    std::string fileName = "[a*.txt";  
  
    if (std::error_code ec = GlobPattern::create(fileName, pattern)) {  
        errorFileName = fileName;  
        return ec;  
    }  
  
    result = pattern.match("...");  
    return std::error_code();  
}
```

Example . clever

```
std::error_code simpleExample(bool &result, std::string *&errorFileName) {
    GlobPattern pattern;
    std::string fileName = "[a*.txt";

    if (std::error_code ec = GlobPattern::create(fileName, pattern)) {
        errorFileName = new std::string(fileName);
        return ec;
    }

    result = pattern.match("...");
    return std::error_code();
}
```

Example . modified

```
int main() {
    bool res;
    std::string *errorFileName = nullptr; // heap alloc in error case

    if (std::error_code ec = simpleExample(res, errorFileName)) {
        std::cerr << "[simpleExample Error] " << getErrorDescription(ec) << " ";
        std::cerr << *errorFileName << "\n";
        delete errorFileName;
        return 0;
    }

    // ... more code ...
    return 0;
}
```

Example . before

```
bool simpleExample() {
    std::string fileName = "[a*.txt";
    Expected<GlobPattern> pattern = GlobPattern::create(std::move(fileName));

    if (auto err = pattern.takeError()) {
        logAllUnhandledErrors(std::move(err), std::cerr, "[Glob Error] ");
        return false;
    }

    return pattern->match("...");
}
```

Example . after

```
Expected<bool> simpleExample() {  
    std::string fileName = "[a*.txt";  
    Expected<GlobPattern> pattern = GlobPattern::create(std::move(fileName));  
  
    if (!pattern)  
        return pattern.takeError();  
  
    return pattern->match("...");  
}
```


Example . after

```
int main() {
    Expected<bool> res = simpleExample();
    if (auto err = res.takeError()) {
        logAllUnhandledErrors(std::move(err), errs(), "[simpleExample Error] ");
        return 0;
    }

    // ... more code ...
    return 0;
}
```

Example . before

```
int main() {  
    if (                simpleExample()) {  
        // ... more code ...  
  
    }  
  
    return 0;  
}
```

Performance

- Concerned about NRVO when seeing code like this?

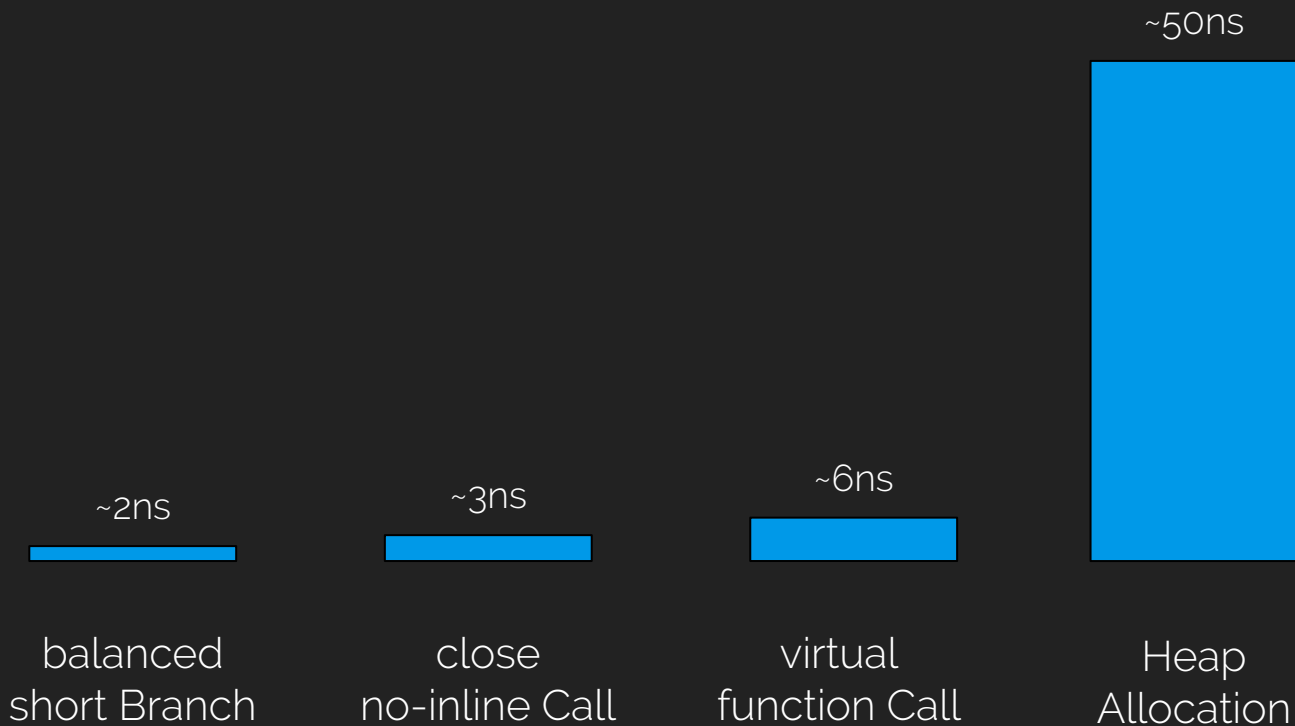
```
return std::move(error);
```

- Concerned about returning polymorphic objects?

Instead of `bool`, `int`, `nullptr`, `std::error_code`

Yes, or course! We only pay for what we get!

Expected overhead category?



Minimal example . std::error_code

```
__attribute__((noinline))
static std::error_code Minimal_ErrorCode(int successRate, int &res) noexcept {
    if (fastrand() % 100 > successRate)
        return std::error_code(9, std::system_category());

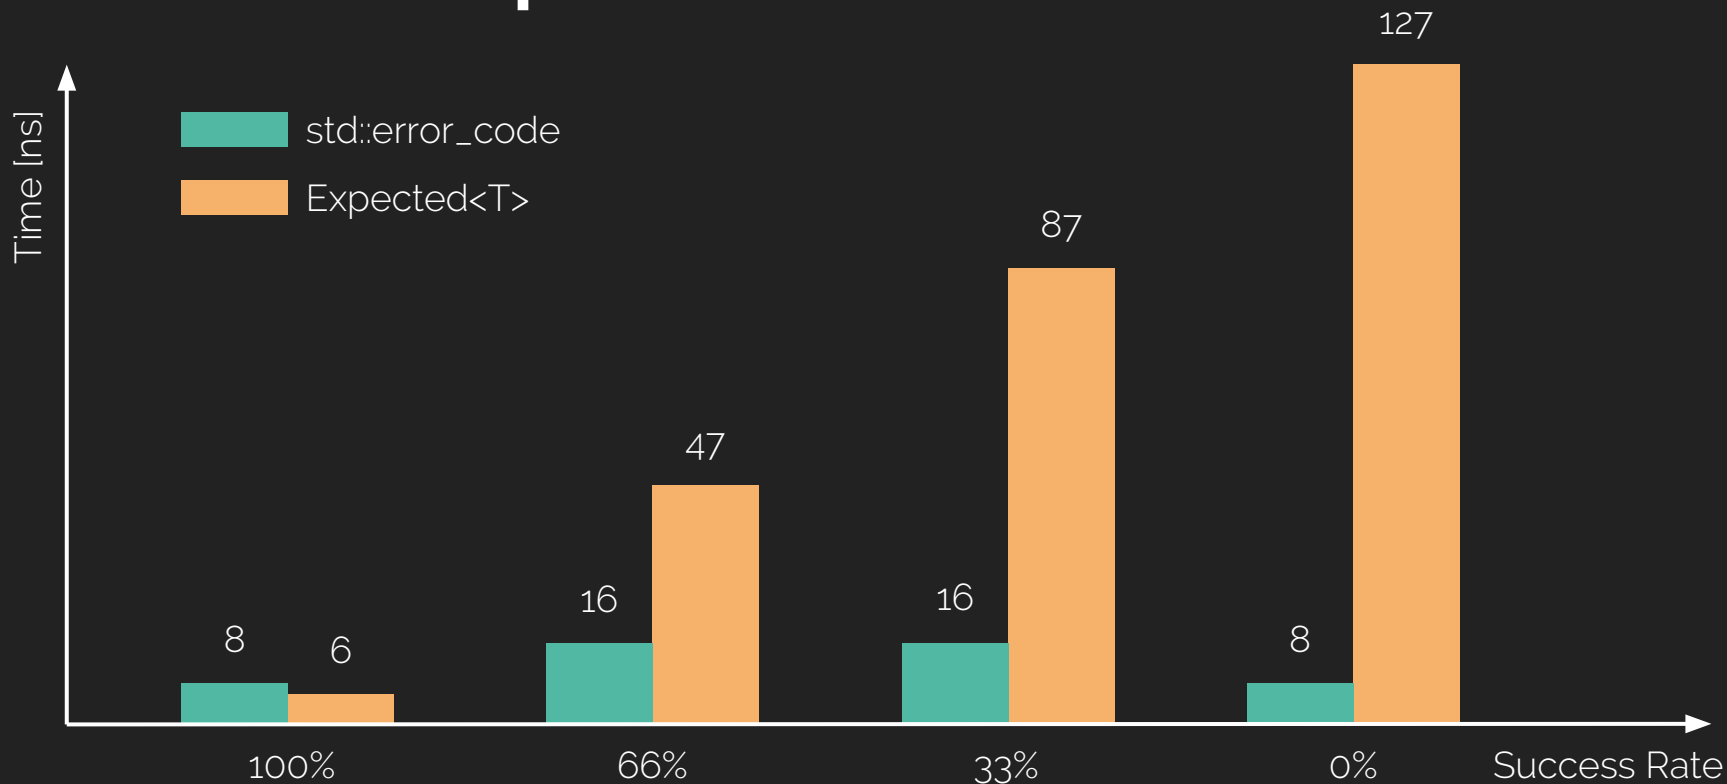
    res = successRate;
    return std::error_code();
}
```

Minimal example . Expected<T>

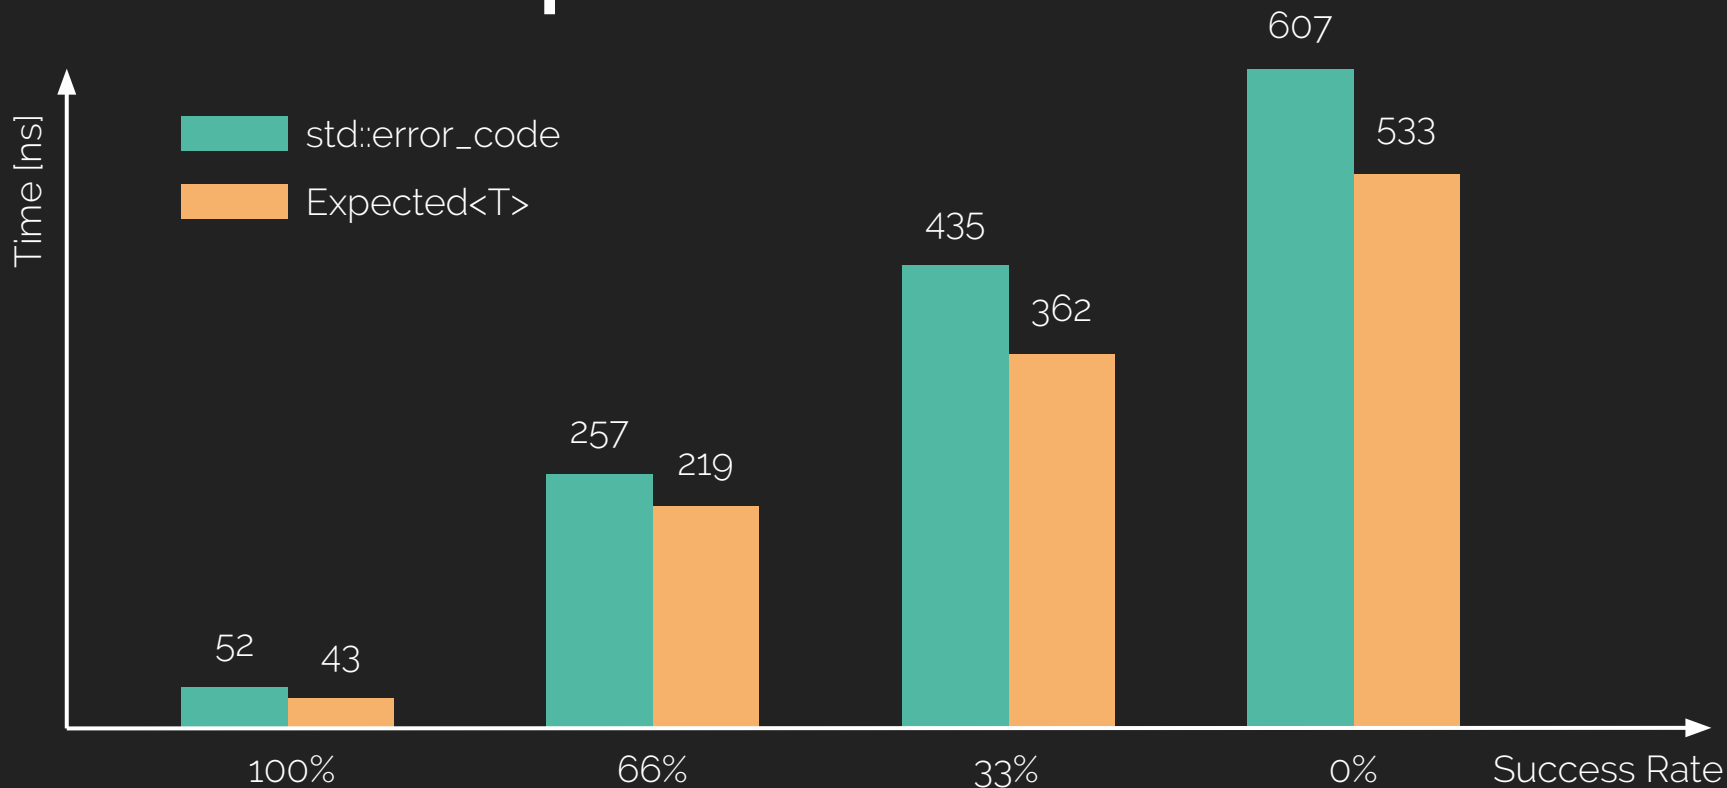
```
__attribute__((noinline))
static llvm::Expected<int> Minimal_Expected(int successRate) noexcept {
    if (fastrand() % 100 > successRate)
        return llvm::make_error<llvm::StringError>(
            "Error Message", llvm::inconvertibleErrorCode());

    return successRate;
}
```

Minimal example



Previous example . after



Expected<T> vs. error code

- ✓ avoid vulnerabilities due to missed errors
- ✓ arbitrarily detailed error descriptions
- ✓ easily propagate errors up the stack
- ✓ no performance loss in success case

Differentiation

Alexandrescu's proposed `Expected<T>`

- made for interop with Exceptions (won't compile with `-fno-exceptions`)
- may pull in implementation-dependent trouble:

```
typedef /*unspecified*/ exception_ptr;
```
- supports `Expected<void>` where LLVM has `Error`

Differentiation

boost::outcome / std::experimental::expected

- interop with exceptions or error codes
- expected<T, E> has error type as template parameter
 - hard to build handy utilities around it
 - IMHO same mistake as static exception specifiers
 - bad versionability, bad scalability: <http://www.artima.com/intv/handcuffsP.html>
- in progress, currently v2, maybe C++20

llvm::Expected<T> vs. others

- ✓ works in real code today
- ✓ supports error concatenation
- ✓ supports error type hierarchies
- ✓ great interop with std::error_code for converting APIs
- ✓ easy to understand, no unnecessary complexity
- not header-only

Test Idea

- Run a piece of code
- Count the number N of valid `Expected<T>` instances
- Execute the code $i = 1 \dots N$ times
- Turn the i 'th valid instance into an error instance
- Each error path will be executed
- Potential issues show up
- Consider running with AddressSanitizer etc.

Dump Example

```
Expected<bool> simpleExample() {  
    std::string fileName = "[a*.txt";  
    Expected<GlobPattern> pattern = GlobPattern::create(std::move(fileName));  
  
    if (pattern) // success case, frequently taken, good coverage  
        return pattern->match("...");  
  
    int x = *(int*)0; // runtime error, unlikely to show up in regular tests  
    return pattern.takeError();  
}
```

Naive Implementation

```
#ifndef NDEBUG
template <typename OtherT>
Expected(OtherT &&Val, typename std::enable_if<...>::type * = nullptr)
    : HasError(false), Unchecked(true)
{
    if (ForceAllErrors::TurnInstanceIntoError()) {
        HasError = true;
        new (getErrorStorage()) error_type(ForceAllErrors::mockError());
        return;
    }

    new (getStorage()) storage_type(std::forward<OtherT>(Val));
}
#else
...
```

Naive Testing

```
int breakInstance = 1..N;
ForceAllErrorsInScope FAE(breakInstance);

Expected<bool> expected = simpleExample();
EXPECT_FALSE(isInSuccessState(expected));

bool success = false;
handleAllErrors(expected.takeError(),
    [&](const ErrorInfoBase &err) { // no specific type information!
        success = true;
    });

EXPECT_TRUE(success);
```


Towards an Error Sanitizer

- Mock correct error type
 - extra info from static analysis → hack Clang
 - runtime support → extend & link LLVM Compiler-RT
- Support cascading errors
 - if error causes more errors, rerun and break all these too
- Avoid breaking instances multiple times
 - deduplicate according to `__FILE__` and `__LINE__`

Towards an Error Sanitizer

→ Biggest challenge:

Missed side effects can cause false-positive results

```
static int SideEffectValue = 0;

llvm::Expected<int> SideEffectExample(bool returnInt) {
    if (returnInt)
        return 0; // ESan breaks the instance created here

    SideEffectValue = 1; // regular errors include this side effect
    return llvm::make_error<CustomError>("Message");
}
```

Towards an Error Sanitizer

- Opinions welcome!
- More news maybe next year

Thks! Questions?

LLVM Programmer's Manual

<http://llvm.org/docs/ProgrammersManual.html#recoverable-errors>

Stripped-down Version of LLVM

<https://github.com/weliveindetail/llvm-expected>

Series of Blog Posts

<http://weliveindetail.github.io/blog/>

Naive Testing Implementation

<https://github.com/weliveindetail/llvm-ForceAllErrors>