
ThinLTO Summaries for Incremental JIT Compilation

Stefan Gränitz

Freelance Dev C++ / Compilers / Audio

LLVM Social Berlin

31. Mai 2018

LTO in a Nutshell

- LTO = Link Time Opt. = Whole Program Opt.
= Cross-Translation-Unit Optimization = Monolithic LTO
- **Compilers** can only inline code that they see, e.g.:
C++ definitions from included headers, but not from cpp's
- **Linkers** can have the full picture, but traditional LTO is expensive and hard to scale

LTO in a Nutshell

- LTO = Link Time Opt. = Whole Program Opt.
= Cross-Translation-Unit Optimization = Monolithic LTO
- **Compilers** can only inline code that they see, e.g.:
C++ definitions from included headers, but not from cpp's
- **Linkers** can have the full picture
expensive and hard to scale

Requires IR code
in object files

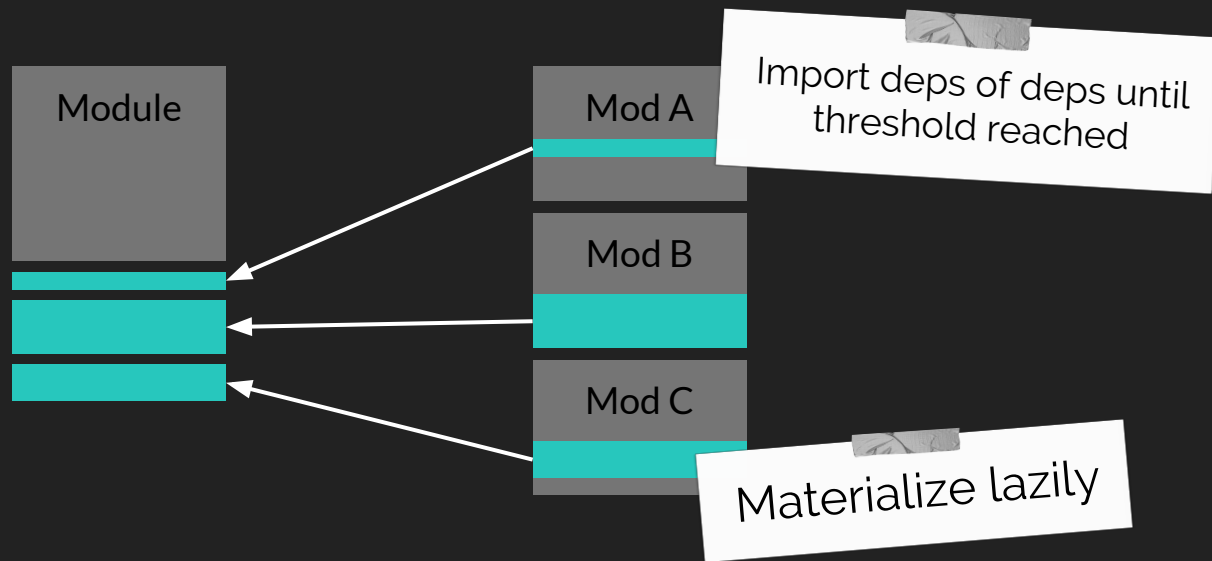
<https://llvm.org/docs/BitCodeFormat.html>
#native-object-file-wrapper-format

ThinLTO in a Nutshell

- **Compiler** emits summaries with call graph info to object files
- **Linker** reads all summaries and builds **global combined index** for Interprocedural Analysis (Thin-Link) based on GUIDs
- Then runs ThinLTO Backend per module:
 - **Import function definitions** from other modules
 - Execute Interprocedural Transformations



ThinLTO Backends

- 1. Import functions + dependencies from other modules




ThinLTO Backends

→ 2. Run inlining and IPO



ThinLTO Backends

→ 3. Drop imports



ThinLTO Benefits

- Small overhead: compact summaries, fast combining
- Good scalability: actual transformations run in parallel
- **Incremental**: simple object file caching + your linker's magic
- Speedup: close to Monolithic LTO
- Development: less manual optimization **avoids discussions on where to put a definition** and can reduce review times dramatically :-)

ThinLTO Command Line

→ Compile:

```
$ clang -flto=thin -O2 file1.c file2.c -c
```

→ Link:

```
$ clang -flto=thin -O2 file1.o file2.o -o a.out
```

→ All in one:

```
$ clang -flto=thin -O2 file1.c file2.c -o a.out
```

More Info

ThinLTO: Scalable and Incremental Link-Time Optimization

Teresa Johnson @ CppCon 2017

<https://www.youtube.com/watch?v=p9nH2vZ2mNo>

ThinLTO: Scalable and Incremental LTO

Teresa Johnson, LLVM Project Blog 2016

<http://blog.llvm.org/2016/06/thinlto-scalable-and-incremental-lto.html>

Clang ThinLTO Documentation

<https://clang.llvm.org/docs/ThinLTO.html>

Incremental Compilation in a Nutshell

- Rebuild only what's necessary
- Terminology:
 - Incremental Linking \approx translation unit granularity
 - Incremental Compilation = higher gran. / smaller entities
- Complexity in languages with C-like translation units:
Incremental Linking \ll Incremental Compilation

Incremental Compilation in a Nutshell

- Rebuild only what's necessary
- Terminology:
 - Incremental Linking \approx translation unit granularity
 - Incremental Compilation = higher gran. / smaller entities

→ Compared with C-like

Incremental

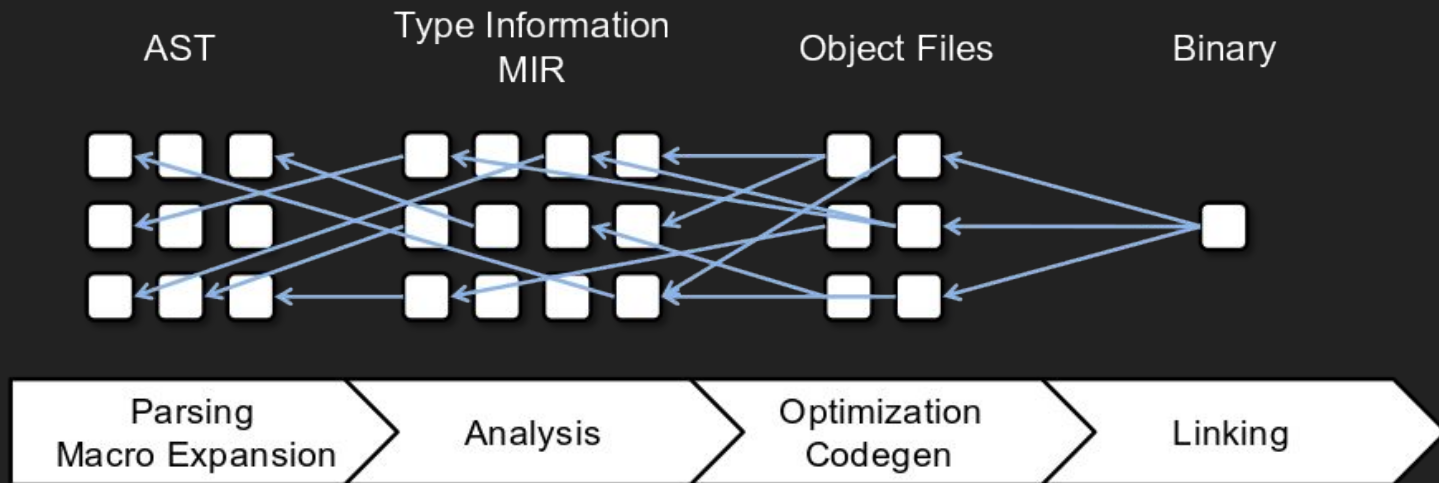
Dynamic linking
to \downarrow turnaround
times! e.g. libLLVM

Incremental compilation

#Entities $\uparrow \Rightarrow$ Overhead \uparrow
 \Rightarrow Eats up benefits

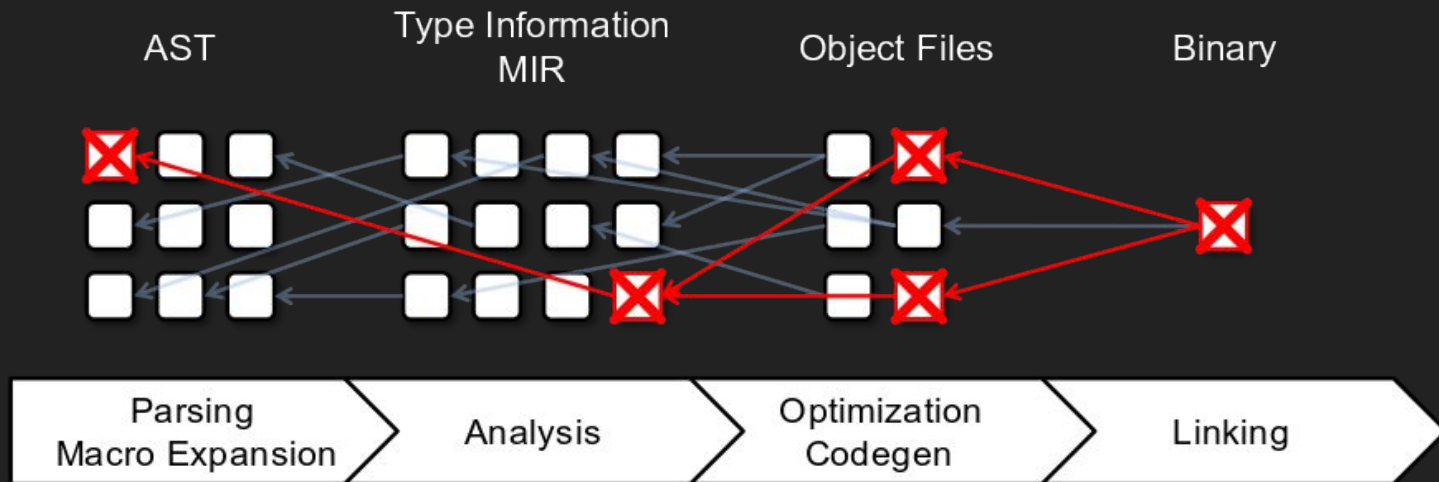
Incremental Compilation in the **rustc** Compiler

→ AST node granularity



Incremental Compilation in the **rustc** Compiler

→ AST node granularity gets complex quickly



Incremental Compilation in the **rustc** Compiler

- AST node granularity gets complex quickly
- Compile-time savings hard to judge: “Up to 15% or More” *
- Conceptually promising, kind-of-academic research

More Info

Incremental Compilation

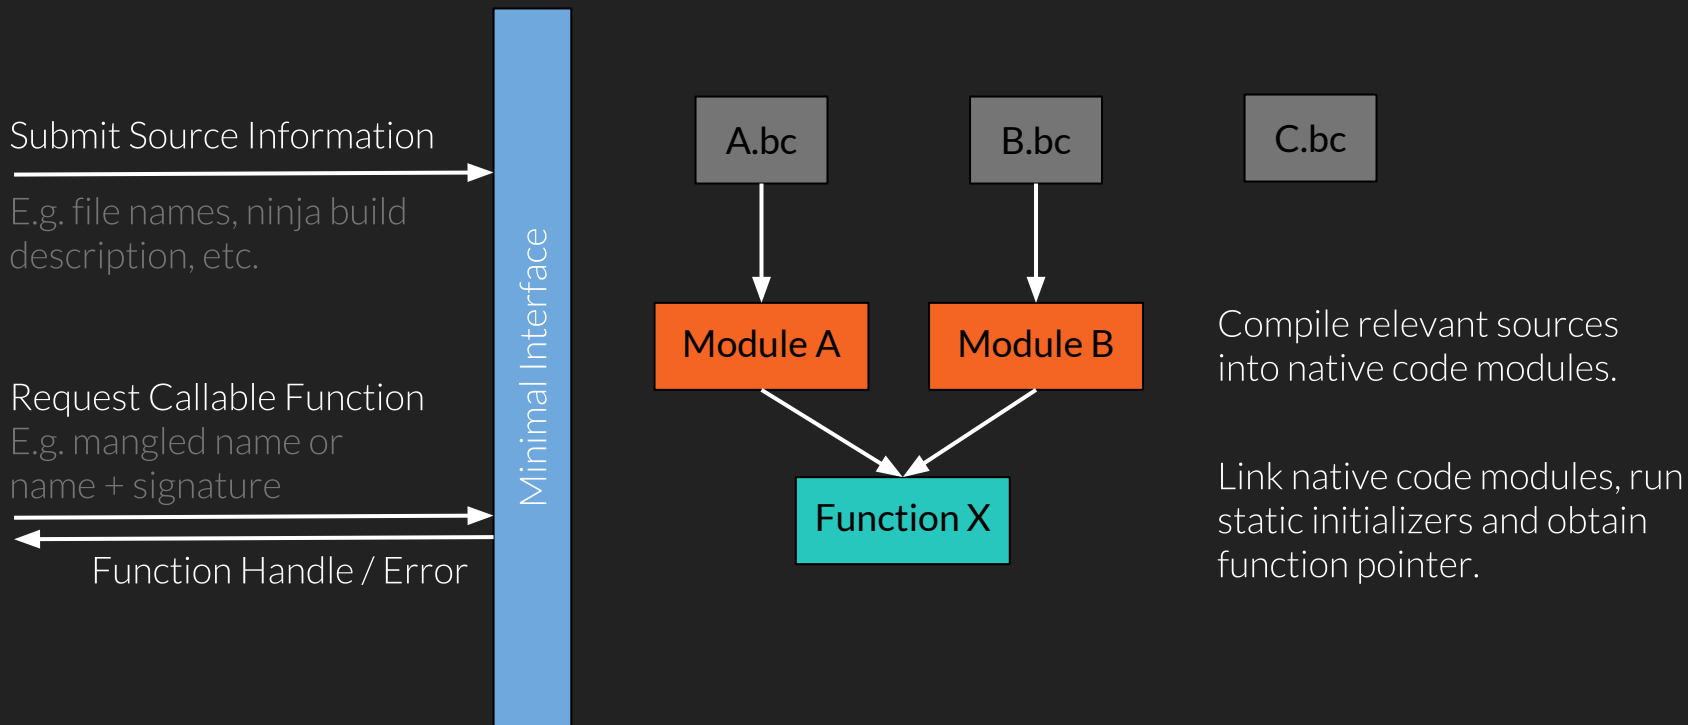
Michael Woerister, Rust Programming Language Blog 2016

<https://blog.rust-lang.org/2016/09/08/incremental.html>

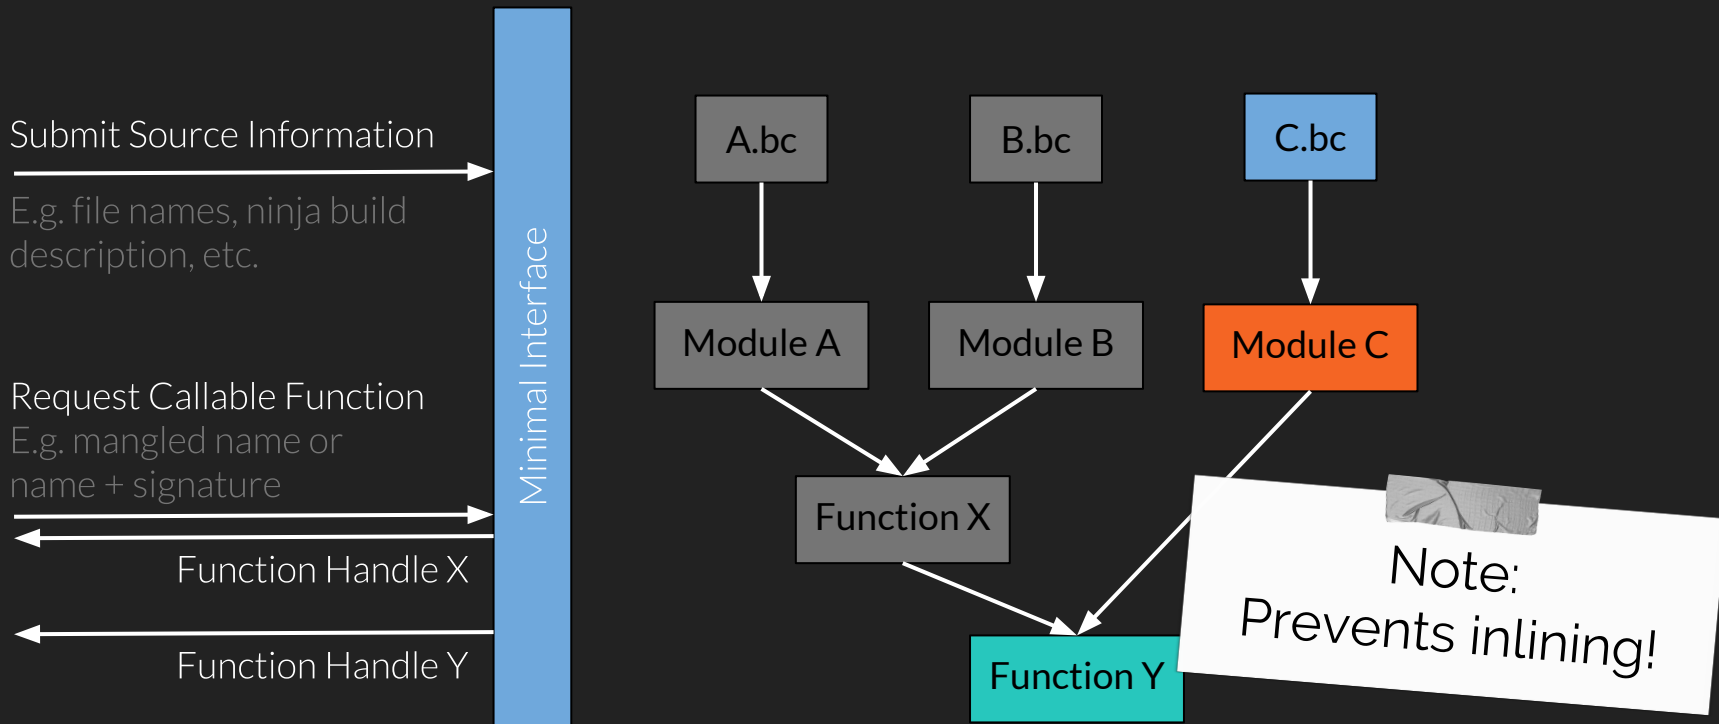
JIT Compilation

- JIT Compiler does both, compile and link
 - ↪ Source representation is LLVM IR
 - ↪ Incremental Compilation \approx Incremental Linking
- No simple separation

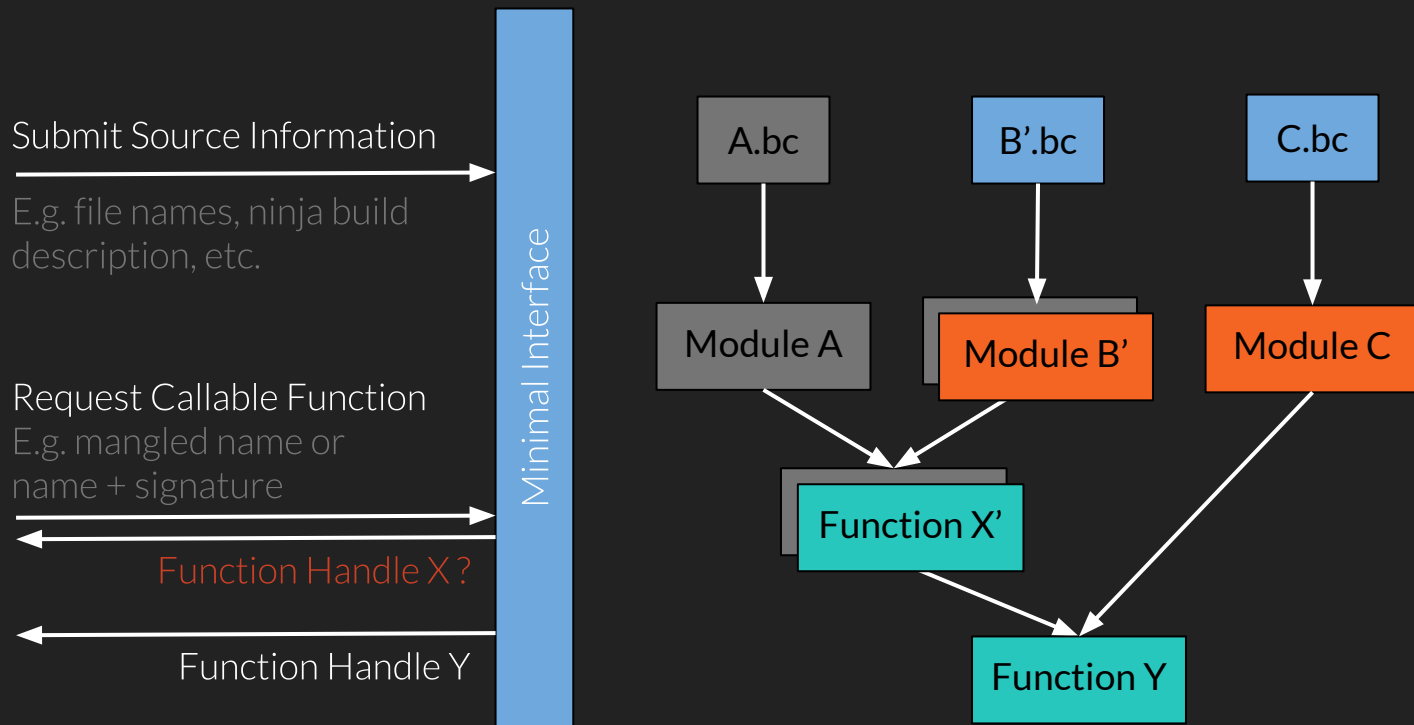
“Monolithic” JIT Compilation



Incremental JIT Compilation



Incremental Just-in-Time Compilation



Incremental JIT Compilation

ThinLTO
Summaries help

→ Requires quite some housekeeping

↪ Track file changes on disk

Source Module Hashes

↪ Determine source dependencies for functions

Call Graph

↪ Decide what needs to be recompiled

↪ Link against existing code

↪ Determine static initializers to rerun

↪ Invalidate existing function handles (and rebind?)

Incremental JIT Compilation

- Requires quite some housekeeping
- More importantly we also **run the code!**
 - ↪ Can we get **Stateful** Incremental JIT Compilation?

Stateful Incremental JIT Compilation

ThinLTO
Summaries help

→ Requires even more housekeeping

↪ Track file changes on disk

Source Module Hashes

↪ Determine source dependencies for functions

Call Graph

↪ Decide what needs to be recompiled

↪ Link against existing code


Name Promotion

↪ Rename duplicate globals and resolve to existing ones

↪ Invalidate existing function handles (and rebind?)

Name Promotion in ThinLTO Backends

- Names of imported globals must not clash with local names!



Name Promotion in ThinLTO Backends

main.cpp

```
int readGlobal();  
void writeGlobal(int x);  
  
static int GlobalVar = 10;  
  
int main() {  
    writeGlobal(1);  
    return readGlobal() + GlobalVar;  
}
```

```
$ clang main.cpp global.cpp -o test  
$ ./test  
$ echo $?  
11
```

Name collision on import



global.cpp

```
static int GlobalVar = 0;  
int readGlobal() { return GlobalVar; }  
void writeGlobal(int x) { GlobalVar = x; }
```

Name Promotion in ThinLTO Backends

Imported pseudo-code main.cpp

```
int readGlobal() { return GlobalVar.llvm.98887675656456543548795; }  
void writeGlobal(int x) { GlobalVar.llvm.98887675656456543548795 = x; }  
  
static int GlobalVar = 10;  
extern int GlobalVar.llvm.98887675656456543548795;  
  
int main() {  
    writeGlobal(1);  
    return readGlobal() + GlobalVar;  
}
```

Postfix with source
module hash

- Like the regular linker, JIT can use hash to find module with the symbol definition

Stateful Incremental Builds

global.cpp

```
#include <stdio>

static int GlobalVar = 0;
int readGlobal() { return GlobalVar; }

void writeGlobal(int x) {
    printf("Modified GlobalVar");
    GlobalVar = x;
}
```

➔ Modification changes module hash

Stateful Incremental Builds

Adjusted pseudo-code global.cpp

```
#include <stdio>

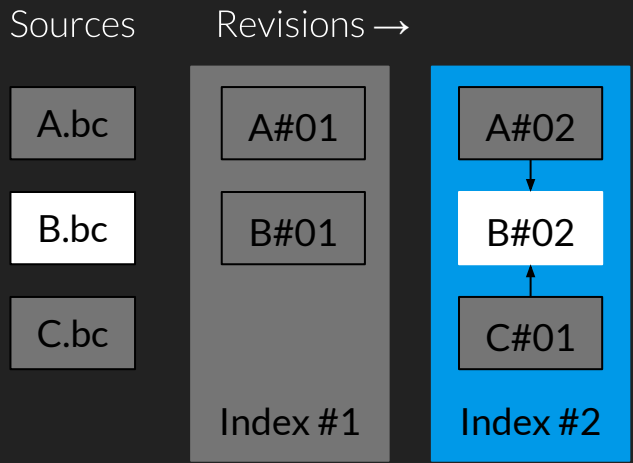
extern int GlobalVar.llvm.98887675656456543548795;
int readGlobal() { return GlobalVar.llvm.98887675656456543548795; }

void writeGlobal(int x) {
    printf("Modified GlobalVar");
    GlobalVar.llvm.98887675656456543548795 = x;
}
```

- ➔ Manually rename own global, so it can be resolved to existing instance in previous revision

Stateful Incremental Builds

- ThinLTO tools tailored for ahead-of-time compilation, e.g. `ModuleSummaryIndex` uses module ID as key and not Hash



- Cannot (easily) keep one Index over multiple revisions.
- [Diff summaries](#) to obtain on-disk changes during Thin-Link
- [Promote stateful globals manually](#) in Thin Backends (lots of extra maintenance!)

Stateful Incremental Builds

- If `ModuleSummaryIndex` keys were hashes, we could easily import from previous revisions:


Sources

A.bc

B.bc

C.bc

Revisions →



Stateful Incremental Builds

- Other ThinLTO implementations used:
 - ↪ `llvm::getLazyBitcodeModule()`
 - ↪ `llvm::readModuleSummaryIndex()`
 - ↪ `llvm::computeDeadSymbols()`
 - ↪ `llvm::ComputeCrossModuleImportForModule()`

Overlap: ThinLTO \leftrightarrow (Stateful) Incremental JIT

- Require global view of the program
- Computationally expensive & huge scale
- Run in two stages
 - ↪ Global analysis (sequential)
 - ↪ Per-module recompilation (in parallel)
- Focus on optimizations in the compiled program
 - ↪ JIT: keep acceptable runtime with high granularity!

Short-Term Goals

- Hack Clang to emit Hashes in ModuleSummaryIndex.
- Build infrastructure upon ThinLTO implementations that directly targets stateful incremental recompilation to get rid of hacky workarounds.
- Extend management infrastructure, e.g. handle cases like function signature changes.
- Load ninja.build descriptions instead of single files.

Long-Term Ideas

- Git-like revision management for program state
 - ↪ Fast-forward merge side-effect-free changes
 - ↪ Let developer resolve conflicts
 - ↪ Allow saving and restoring state snapshots
 - ↪ Fall back to restart on fail

Long-Term Ideas

- Git-like revision management for program state
- Emit program summary data to executable (or extra file)
 - ↪ Like debug information for program state

Long-Term Ideas

- Git-like revision management for program state
- Emit program summary data to executable (or extra file)
 - ↪ Like debug information for program state
- Package JIT as dylib, so it can be used like a sanitizer

Long-Term Ideas

- Git-like revision management for program state
- Emit program summary data to executable (or extra file)
 - ↪ Like debug information for program state
- Package JIT as dylib, so it can be used like a sanitizer
- Implement support in common developer tools (e.g. LLDB) to modify executables based on program summary data

Thanks!

Let's dive into the code!

<https://github.com/weliveindetail/ThinLtoJit>